

incomplete rollout segment ($T=512$, no terminal state), $\kappa = 0.95$

last $(1-\kappa)T$ steps zeroed

policy
gradient

advantages kept ($\lfloor \kappa T \rfloor$ steps)

value
loss

value targets trained on *all* T steps

0

$\lfloor \kappa T \rfloor$ T

timestep t

